

## Laboratório 6 - Treinar um modelo neural de tradução



### Diário de bordo

Código:

<https://colab.research.google.com/drive/13YMtt4IrbYL9T3SoVbXEG0MrqHiPnRBx?usp=sharing>

## 1 Hyperparâmetros:

source\_embedding\_size=24  
target\_embedding\_size=24  
encoding\_size=32  
batch\_size=32  
num\_epochs=100  
learning\_rate=0.0005



```
(encoder): NMTEncoder(  
  (source_embedding): Embedding(3025, 24, padding_idx=0)  
  (birnn): GRU(24, 32, batch_first=True, bidirectional=True)  
(decoder): NMTDecoder(  
  (target_embedding): Embedding(4911, 24, padding_idx=0)  
  (gru_cell): GRUCell(88, 64)  
  (hidden_map): Linear(in_features=64, out_features=64, bias=True)  
  (classifier): Linear(in_features=128, out_features=4911, bias=True)
```

[INFO] epoch 22, **train loss 1.3188232963545286, val loss 2.5745678397475698**

[INFO] **train\_acc 70.17609968704853, val acc 61.43141399014763**

[INFO] best validation loss updated and checkpoint saved

### Evaluate NMT



[INFO] **mean result: 0.4276705749353432, median result: 0.4061551592130281**

### Comentários:

Overfitting, vamos diminuir a taxa de aprendizagem.

## 2 Hyperparâmetros:

source\_embedding\_size=24  
target\_embedding\_size=24  
encoding\_size=32  
batch\_size=32  
num\_epochs=50  
learning\_rate=0.00005



[INFO] using device cuda

```
NMTModel(  
  (encoder): NMTEncoder(  
    (source_embedding): Embedding(3025, 24, padding_idx=0)  
    (birnn): GRU(24, 32, batch_first=True, bidirectional=True)  
  )  
  (decoder): NMTDecoder(  
    (target_embedding): Embedding(4911, 24, padding_idx=0)  
    (gru_cell): GRUCell(88, 64)  
    (hidden_map): Linear(in_features=64, out_features=64, bias=True)  
    (classifier): Linear(in_features=128, out_features=4911, bias=True)  
  )  
)
```

[INFO] epoch 49, **train loss 2.9866147559985787, val loss 3.5085229209212003**

[INFO] **train\_acc 50.221524468939506, val acc 50.350192252949405**

[INFO] best validation loss updated and checkpoint saved

### Evaluate NMT



[INFO] **mean result: 0.28632547839254285, median result: 0.2703810696954543**

### Comentários:

Os resultado pioraram em relação ao 1º treinamento.

### 3 Hyperparâmetros:

source\_embedding\_size=64  
target\_embedding\_size=64  
encoding\_size=32  
batch\_size=32  
num\_epochs=100  
learning\_rate=0.0005



[INFO] using device cuda

```
NMTModel(  
  (encoder): NMTEncoder(  
    (source_embedding): Embedding(3025, 64, padding_idx=0)  
    (birnn): GRU(64, 32, batch_first=True, bidirectional=True)  
  )  
  (decoder): NMTDecoder(  
    (target_embedding): Embedding(4911, 64, padding_idx=0)  
    (gru_cell): GRUCell(128, 64)  
    (hidden_map): Linear(in_features=64, out_features=64, bias=True)  
    (classifier): Linear(in_features=128, out_features=4911, bias=True)  
  )  
)
```

[INFO] epoch 19, **train loss 1.4317084023827007, val loss 2.6423677929112173**

[INFO] **train\_acc 67.99527847189836, val acc 59.980514102363784**

[INFO] best validation loss updated and checkpoint saved

#### Evaluate NMT



[INFO] **mean result: 0.359568668030878, median result: 0.3328359601570994**

#### Comentários:

Melhor resultado até aqui, contudo temos muito overfitting. Vamos aumentar a taxa de aprendizagem.

### 4 Hyperparâmetros:

source\_embedding\_size=64  
target\_embedding\_size=64  
encoding\_size=32  
batch\_size=32  
num\_epochs=100  
learning\_rate=0.005



[INFO] using device cuda

```
NMTModel(  
  (encoder): NMTEncoder(  
    (source_embedding): Embedding(3025, 64, padding_idx=0)  
    (birnn): GRU(64, 32, batch_first=True, bidirectional=True)  
  )  
  (decoder): NMTDecoder(  
    (target_embedding): Embedding(4911, 64, padding_idx=0)  
    (gru_cell): GRUCell(128, 64)  
    (hidden_map): Linear(in_features=64, out_features=64, bias=True)  
    (classifier): Linear(in_features=128, out_features=4911, bias=True)  
  )  
)
```

[INFO] epoch 6, **train loss 0.5302001076832155, val loss 2.0058176751996646**

[INFO] **train\_acc 84.79429008294372, val acc 69.7204647070264**

[INFO] best validation loss updated and checkpoint saved

#### Evaluate NMT



[INFO] **mean result: 0.5454695553108071, median result: 0.5439094815855025**

#### Comentários:

Melhor resultado alcançado, contudo o overfitting não diminuiu. Habilitaremos a CrossEntropyLoss.

## 5 Hyperparâmetros:

source\_embedding\_size=64  
target\_embedding\_size=64  
encoding\_size=32  
batch\_size=32  
num\_epochs=100  
learning\_rate=0.005  
loss\_func =  
**nn.CrossEntropyLoss()**





### Comentários:

Paramos o treinamneto na época 18, pois não observamos nenhuma melhoria no overfitting. Vamos habilitar o method 2: sample from distribution

## 6 Hyperparâmetros:

source\_embedding\_size=64  
target\_embedding\_size=64  
encoding\_size=32  
batch\_size=32  
num\_epochs=100  
learning\_rate=0.005  
loss\_func = nn.CrossEntropyLoss()  
**# method 2: sample from distribution**  
**candidate\_input =**  
**torch.multinomial(p\_y\_t\_index,**  
**1).squeeze()**




### Comentários:

Paramos o treinamneto na época 18, pois não observamos nenhuma melhoria no overfitting. Vamos habilitar a regularização L2.

## 7 Hyperparâmetros:

source\_embedding\_size=64  
target\_embedding\_size=64  
encoding\_size=32  
batch\_size=32  
num\_epochs=100  
learning\_rate=0.005  
l2\_regularization = **0.001**



```
[INFO] using device cpu
NMTModel(
  (encoder): NMTEncoder(
    (source_embedding): Embedding(3025, 64, padding_idx=0)
    (birnn): GRU(64, 32, batch_first=True, bidirectional=True)
  )
  (decoder): NMTDecoder(
    (target_embedding): Embedding(4911, 64, padding_idx=0)
    (gru_cell): GRUCell(128, 64)
    (hidden_map): Linear(in_features=64, out_features=64, bias=True)
    (classifier): Linear(in_features=128, out_features=4911, bias=True)
  )
)
[INFO] epoch 45, train loss 3.076054169002329, val loss 2.8562438409836557
[INFO] train_acc 54.99467830925427, val acc 54.526114394602814
[INFO] best validation loss updated and checkpoint saved
```


#### Evaluate NMT

[INFO] **mean** result: **0.331**15399770341614, **median** result: **0.321**9038573439098

#### Comentários:

Diminuimos o overfitting, mas a acurácia diminui expressivamente.

- 8 **Hyperparâmetros:**  
 source\_embedding\_size=64  
 target\_embedding\_size=64  
 encoding\_size=32  
 batch\_size=32  
 num\_epochs=40  
 learning\_rate=**0.0005**  
 l2\_regularization = 0.001



```
[INFO] using device cpu
NMTModel(
  (encoder): NMTEncoder(
    (source_embedding): Embedding(3025, 64, padding_idx=0)
    (birnn): GRU(64, 32, batch_first=True, bidirectional=True))
  (decoder): NMTDecoder(
    (target_embedding): Embedding(4911, 64, padding_idx=0)
    (gru_cell): GRUCell(128, 64)
    (hidden_map): Linear(in_features=64, out_features=64, bias=True)
    (classifier): Linear(in_features=128, out_features=4911, bias=True)))
[INFO] epoch 39, train loss 3.231030480066935, val loss 2.9168540337046633
[INFO] train_acc 54.33450864742833, val acc 54.39915909182199
[INFO] best validation loss updated and checkpoint saved
```

#### Evaluate NMT

[INFO] **mean** result: **0.321**770244052043, **median** result: **0.306**3079764817904



#### Comentários:

Diminuimos o overfitting, mas a acurácia não melhorou.

## 9 Hyperparâmetros:

source\_embedding\_size=64  
target\_embedding\_size=64  
encoding\_size=32  
batch\_size=32  
num\_epochs=40  
learning\_rate=0.0005  
l2\_regularization = 0.001

```
loss_func = nn.CrossEntropyLoss()
optimizer =
optim.Adagrad(model.parameters(),
               lr=learning_rate,
               weight_decay=l2_regularization)
scheduler = optim.lr_scheduler.ReduceLROnPlateau
(optimizer=optimizer, mode='min', factor=0.5, patience=1)
```



```
[INFO] using device cpu
NMTModel(
  (encoder): NMTEncoder(
    (source_embedding): Embedding(3025, 64, padding_idx=0)
    (birnn): GRU(64, 32, batch_first=True, bidirectional=True)
  )
  (decoder): NMTDecoder(
    (target_embedding): Embedding(4911, 64, padding_idx=0)
    (gru_cell): GRUCell(128, 64)
    (hidden_map): Linear(in_features=64, out_features=64, bias=True)
    (classifier): Linear(in_features=128, out_features=4911, bias=True)))
```

[INFO] epoch 0, train **loss 509.4013233920982**, val **loss 6.958291835472233**

[INFO] train\_acc **22.09021296112598**, val acc **26.224576409561**

[INFO] best validation loss updated and checkpoint saved

[INFO] epoch 39, train loss **434.5768570749383**, val **loss 4.123427863980903**

[INFO] train\_acc **37.94451676227942**, val acc **38.15299286013642**

[INFO] best validation loss updated and checkpoint saved



[INFO] mean result: **0.08892326948954755**, median result: **0.0439975593063474**

**Comentário:** Foi mudado o otimizador ADAM por Adagrad, mas a perda no treinamento não caiu, e a acurácia também não melhorou, vamos voltar para o ADAM e diminuir o tamanho do embedding\_size

## 10 Hyperparâmetros:

source\_embedding\_size=24  
target\_embedding\_size=24  
encoding\_size=32  
batch\_size=32  
num\_epochs=40  
learning\_rate=0.0005  
l2\_regularization = 0.001

```
loss_func = nn.CrossEntropyLoss()
optimizer =
optim.Adam(model.parameters(),
            lr=learning_rate,
            weight_decay=l2_regularization)
scheduler = optim.lr_scheduler.ReduceLROnPlateau
(optimizer=optimizer, mode='min', factor=0.5, patience=1)
```



```
[INFO] using device cpu
NMTModel(
  (encoder): NMTEncoder(
    (source_embedding): Embedding(3025, 24, padding_idx=0)
```

```

(birnn): GRU(24, 32, batch_first=True, bidirectional=True)
)
(decoder): NMTDecoder(
  (target_embedding): Embedding(4911, 24, padding_idx=0)
  (gru_cell): GRUCell(88, 64)
  (hidden_map): Linear(in_features=64, out_features=64, bias=True)
  (classifier): Linear(in_features=128, out_features=4911, bias=True)
)
)

```

[INFO] epoch 0, train loss **175.9322303102728**, val loss **4.200261608498997**

[INFO] train\_acc **28.806423796941488**, val acc **36.35309480582181**

[INFO] epoch 39, train loss **3.265248627411694**, val loss **3.0413855255627245**

[INFO] train\_acc **52.93673662988597**, val acc **52.47779200615325**

[INFO] mean result: **0.28951629576458243**, median result: **0.28023947440512487**

## 11 Hyperparâmetros:

source\_embedding\_size=24

target\_embedding\_size=24

encoding\_size=32

batch\_size=32

num\_epochs=40

learning\_rate=0.0005

l2\_regularization = **0.0001**

loss\_func = nn.CrossEntropyLoss()



optimizer =

optim.**Adam**(model.parameters(),  
lr=learning\_rate,

weight\_decay=l2\_regularization)

scheduler = optim.lr\_scheduler.ReduceLROnPlateau

(optimizer=optimizer, mode='min', factor=0.5, patience=1)



[INFO] using device cpu

NMTModel(

(encoder): NMTEncoder(

(source\_embedding): Embedding(3025, 24, padding\_idx=0)

(birnn): GRU(24, 32, batch\_first=True, bidirectional=True)

)

(decoder): NMTDecoder(

(target\_embedding): Embedding(4911, 24, padding\_idx=0)

(gru\_cell): GRUCell(88, 64)

(hidden\_map): Linear(in\_features=64, out\_features=64, bias=True)

(classifier): Linear(in\_features=128, out\_features=4911, bias=True)

)

)

[INFO] epoch 0, train loss **22.4345690208569**, val loss **4.12170874486204**

[INFO] train\_acc **27.7706076071854**, val acc **38.98964887069802**

[INFO] epoch 39, train loss **2.012769417595447**, val loss **2.418139670715957**

[INFO] train\_acc **64.64790593242903**, val acc **60.32771530425084**

[INFO] mean result: **0.3684211164448593**, median result: **0.3547357381807338**



**Comentário:** O início da perda de treinamento começou com um bom parâmetro em comparação com os outros testes. Houve uma melhora no resultado final, tanto na perda validação e treinamento, e as acurácias também melhoraram, porém começou a ter overfitting. Vamos diminuir o learning rate

## 12 Hyperparâmetros:

source\_embedding\_size=24  
target\_embedding\_size=24  
encoding\_size=32  
batch\_size=32  
num\_epochs=40  
learning\_rate=**0.00005**  
l2\_regularization = 0.0001

```
loss_func = nn.CrossEntropyLoss()
optimizer =
optim.Adam(model.parameters(),
            lr=learning_rate,
```

```
weight_decay=l2_regularization)
scheduler = optim.lr_scheduler.ReduceLROnPlateau
(optimizer=optimizer, mode='min', factor=0.5, patience=1)
```



[INFO] using device cpu

```
NMTModel(
  (encoder): NMTEncoder(
    (source_embedding): Embedding(3025, 24, padding_idx=0)
    (birnn): GRU(24, 32, batch_first=True, bidirectional=True)
  )
  (decoder): NMTDecoder(
    (target_embedding): Embedding(4911, 24, padding_idx=0)
    (gru_cell): GRUCell(88, 64)
    (hidden_map): Linear(in_features=64, out_features=64, bias=True)
    (classifier): Linear(in_features=128, out_features=4911, bias=True)
  )
)
```



[INFO] epoch 0, train loss **27.216734180115814**, val loss **7.1493658237769955**  
[INFO] train\_acc **11.153796903258474**, val acc **18.10693371622847**  
[INFO] train\_acc **49.53038416246259**, val acc **49.47142778088714**  
[INFO] best validation loss updated and checkpoint saved

[INFO] mean result: **0.2876296118027262**, median result: **0.274651895686728**

**Comentário:** Houve melhora no overfitting porém os resultados finais pioraram, por enquanto o melhor resultado foi o **teste 11**

## 13 Hyperparâmetros:

source\_embedding\_size=24  
target\_embedding\_size=24  
encoding\_size=32  
batch\_size=32  
num\_epochs=40  
learning\_rate=0.00005  
l2\_regularization = **0.00001**



[INFO] using device cpu

```
NMTModel(
  (encoder): NMTEncoder(
    (source_embedding): Embedding(3025, 24, padding_idx=0)
    (birnn): GRU(24, 32, batch_first=True, bidirectional=True)
  )
  (decoder): NMTDecoder(
    (target_embedding): Embedding(4911, 24, padding_idx=0)
    (gru_cell): GRUCell(88, 64)
    (hidden_map): Linear(in_features=64, out_features=64, bias=True)
  )
)
```



```
(classifier): Linear(in_features=128, out_features=4911, bias=True)
)
)
```

[INFO] train\_acc **9.98763488656869**, val acc **26.826147907940516**

[INFO] best validation loss updated and checkpoint saved

[INFO] epoch 39, train loss **3.9877724538769663**, val loss **3.343913543419759**

[INFO] train\_acc **49.83101039829001**, val acc **50.10893205099904**

[INFO] mean result: **0.2835466244526874**, median result: **0.266633273235631**

Comentário: Aumentando o learning rate melhorou os parâmetros de de perda, mas não teve impacto significativo na acurácia. Mais uma vez, o melhor treinamento foi o **teste 11**